

PreCex: Counterexample-Driven Pre-Synthesis RTL Defect Localization and Repair

Integrated full draft (work-in-progress). Author: Toylog | 2026-08-04 Sections assembled from paper/draft/{intro,method,results,related,threats}.md. All numbers trace to corrected BMC-judged data (experiments/runs/).

Abstract

Cross-cycle behavioral defects — register-transfer-level (RTL) bugs where weak simulation passes but formal verification fails — are among the hardest to localize and repair before synthesis. We present PreCex, a fully open-source, counterexample-driven pipeline (iverilog/yosys/SymbiYosys/Z3 + a large language model, LLM) that converts a failing formal trace into structured evidence, localizes the defective region, generates a minimal patch, and re-verifies it under bounded model checking (BMC) with a golden-first criterion.

On a 34-sample L3 benchmark across 6 modules and 7 error classes (408 LLM repair runs), all four evidence representations — raw logs, structured evidence (B), semanticized counterexamples (C), and FVDebug-style causal graphs (D) — reach 100% repair under BMC. Structured evidence achieves the highest localization precision (61.8%), significantly better than raw logs (47.1%, $p=0.0035$) and causal graphs (49.0%, $p=0.0164$), while causal graphs are the cheapest (\$1.56 vs \$2.72, 43% lower). Mutation analysis kills 88.5% of strong and 81.8% of constant mutants; an independent T2 audit passes 408/408 runs with zero interface or assertion-tampering changes. A criterion-reversal audit shows prove/k-induction misjudges 78 correct repairs on gated sequential assertions, motivating our golden-first BMC criterion — a methodological lesson for automated repair evaluation.

Keywords: RTL debugging; formal verification; counterexample; LLM repair; mutation analysis

1. Introduction

Problem

RTL designs are increasingly complex, and cross-cycle behavioral defects (L3: weak testbenches pass but formal verification fails) are among the hardest to localize and repair before synthesis. Existing LLM-based repair pipelines struggle because they either feed raw logs/waveforms (information overload) or rely on reference models (not always available pre-synthesis). Recent work (FVDebug, GROVE) shows the value of structured counterexample understanding, but lacks: (1) an open-source end-to-end pipeline, (2) cross-cycle bug-specific study, (3) a rigorous verification-sufficiency loop, and (4) a controlled ablation of evidence representations.

Motivation

We observe that the *evidence representation* given to the LLM is a key but understudied axis. We systematically compare four evidence settings on the same 34-sample L3 dataset (408 LLM repair runs):

- A: raw cex log/VCD
- B: structured evidence (EvidenceEngine JSON)
- C: counterexample semanticization (CexSemantizer: cycle events + state trace + fault cone + NL summary)
- D: FVDebug-style causal graph (deterministic extraction)

Contributions

1. **Open-source counterexample-driven repair pipeline** (iverilog/yosys/sby/z3 + MiniMax M3 API), fully reproducible; 34-sample L3 dataset with golden double-check and sufficiency audits.
2. **Evidence-representation ablation:** all settings reach 100% repair under a consistent BMC criterion; B (structured) significantly improves localization (61.8% vs A 47.1%, $p=0.0035$; vs D 49.0%, $p=0.0164$); D is cheapest (\$1.56); C adds cost without significant gain ($p=0.404$).
3. **Verifier-sufficiency loop:** mutation family (strong 88.5% / const 81.8% killed), vacuity control (delete), deterministic T2 audit (interface/assertion/evidence-loop 306+102 all pass).
4. **Criterion-reversal lesson:** prove/k-induction misjudges 78 correct repairs on gated sequential assertions (golden itself UNKNOWN); BMC criterion + golden-first validation is required — a methodological contribution for reliable automated repair.

2. Method




Figure 1: PreCex pipeline

*Figure 1: PreCex pipeline. A failing formal trace is parsed into structured evidence (B), optionally semanticized (C) or summarized as a causal graph (D), used by the LLM repairer to produce a minimal diff, and re-verified under BMC with a golden-first criterion; failures re-enter the loop.

System architecture

Core 4 components + 2 support elements (per docs/ .md §4):

1. **EvidenceEngine** — parses cex.log into unified JSON (error_type/module/file/line/code_slice/signals/triggers).
2. **CexSemantizer** — converts VCD into cycle-event table + state trace + fault cone + NL summary (window=8 compression).
3. **LocalRepairer** — LLM (MiniMax M3, temperature 0.2) given design + assertions + evidence, outputs locate + minimal unified diff; retries ≤ 2 .
4. **Verifier** — three-pass gate: (1) iverilog compile 0 error, (2) weak testbench sim all-green, (3) sby smtbmc+z3 BMC no counterexample (primary criterion); golden double-check.

Support: Controller (experiment variable), BugBench-PS (dataset assets).

Evidence settings (controlled ablation)

Setting	Evidence	Extraction	Cost/run
A	raw cex.log + VCD head/tail	none	baseline
B	evidence.json	deterministic	low
C	semantics.json (cycle events + state trace + fault cone + NL summary, window=8)	LLM semanticization	high
D	FVDebug-style causal graph (failed assertion + fault cone + full state trace + trigger)	deterministic (zero LLM)	lowest

Same prompt family, only evidence segment replaced (anti-bias protocol).

Repair criterion (revised)

- Primary: BMC (verify.sby), consistent with golden verification (verify_golden.sby).
- prove/k-induction (verify_repair.sby) kept as supplementary reference only — it is **not** a failure criterion, because it does not converge on gated sequential assertions (golden itself UNKNOWN, rc=4; 78 correct repairs were false-FAILED under the old criterion).
- **Golden-first rule:** any repair criterion must first pass on golden (criterion-vs-golden consistency check).

Sufficiency audits

- Strong mutation (comparator inversion, d16): 400 mutants, 88.5% killed.
- Const mutation (parameter constants DATA_W/ADDR_W/DEPTH/DIV/HALF): 484 mutants, 81.8% killed.
- Delete mutation (assertion removal): vacuity control.
- T2 deterministic audit: interface change 0, assertion tampering 0, evidence-loop 0 (306 + 102).

3. Results

RQ1: Can counterexample-grounded evidence enable reliable repair of cross-cycle (L3) RTL defects?

Setup. 34 L3 samples (s04–s37) x 4 evidence settings x 3 seeds = 408 LLM repair runs (A raw log / B structured evidence / C counterexample semanticization / D FVDebug-style causal graph, added 2026-08-04). Repair success judged by BMC (verify.sby) with golden double-check.

Setting	n	Loc Top-1	Repair (BMC)	Tokens	Cost (USD)
A (raw)	102	47.1%	100%	1.99M	\$2.78
B (structured)	102	61.8%	100%	1.76M	\$2.72
C (semanticized)	102	56.9%	100%	3.50M	\$4.17
D (causal graph)	102	49.0%	100%	1.07M	\$1.56




Figure 2: Four-setting localization and cost

Figure 2: *loc_top1* precision (bars) and LLM cost (dashed line) by evidence setting. *B*’s precision lead over *A* and *D* is significant (paired McNemar).

Finding 1. All four settings reach 100% repair under a consistent BMC criterion — evidence representation does not gate *whether* the model can fix, but *how precisely and cheaply*.

Finding 2. B (structured evidence) dominates on precision (61.8%); D (FVDebug-style causal graph) is cheapest (\$1.56, 43% below B; 1.07M tokens) at 49.0% precision - a clear precision/cost

trade-off. C’s semanticization adds 1.7–2× cost with no repair gain (consistent with Gate-0 prestudy).

Significance (paired McNemar on 102 loc outcomes). B vs A: $p=0.0035$ (significant); B vs D: $p=0.0164$ (significant); B vs C: $p=0.404$; C vs D: $p=0.186$. B’s precision lead over raw-log (A) and causal-graph (D) is statistically significant; C’s semanticization adds cost without significant precision gain.

Finding 3 (loc difficulty by error type). Single-point semantic errors are easy to locate (edge 100%, width_trunc 85.2%); cross-state/handshake errors are hard (state_trans 36.1%, handshake 27.8%). The relative difficulty ordering is stable across settings. Setting-level localization by error class (per-setting n: 24/24/24/24 for state_trans, 12/12/12/12 for handshake, 3/3/3/3 for edge):

Error class	A	B	C	D
state_trans	33.3%	37.5%	37.5%	25.0%
handshake	16.7%	33.3%	33.3%	41.7%
reset	55.6%	72.2%	83.3%	44.4%
fifo_full_empty	53.3%	53.3%	60.0%	46.7%
boundary_wrap	42.9%	81.0%	57.1%	57.1%
width_trunc	88.9%	100%	66.7%	100%
edge	100%	100%	100%	100%

Figure 3: *loc_top1* precision by error class and setting. The hard-class ordering (state_trans/handshake below width_trunc/edge) holds across all settings.

B is best or tied-best on boundary_wrap (81.0%), width_trunc (100%), and state_trans (37.5%); D is highest on handshake (41.7%, $n=12$ per setting — small- n caveat). The hard-class ordering (state_trans/handshake below width_trunc/edge) holds in every setting.

Non-LLM baselines. To establish that the LLM’s localization is non-trivial, we evaluate simple heuristic baselines on the same exact-match criterion (candidate line == buggy_inject_line, 34 samples):

Baseline	loc_top1
First assertion line	0.0%
Any assertion line	0.0%
First line mentioning a failing signal	0.0%
Random line (expected value)	0.50%

All four LLM settings (47.1%–61.8%) far exceed these baselines; structured evidence (B) achieves 61.8% — over 120x the random-line expectation. The assertion- and signal-based heuristics never hit because injected defects live in functional/sequential logic, not in the assertion block itself.

RQ2: Is the repair trustworthy? (Verifier sufficiency + patch quality)

Sufficiency family (mutation analysis on golden):




Figure 3: Error-class x setting heatmap

Mutation family	Mutants	Killed	Rate
Strong (comparator inversion)	400	354	88.5%
Const (parameter constant replacement)	484	396	81.8%
Delete (assertion removal)	sampled	all PASS	vacuity control

Module-level: axi/fifo strong=100%; uart_rx weakest (strong 55.0%, const 40.9%) — structural assertions are insensitive to baud-timing constant shifts (documented limitation, future work).

T2 audit (deterministic, 408/408 pass: A/B/C 306 + D 102): interface changes 0, assertion tampering 0, evidence-loop failures 0; loc_dev median = 0 (62.7% exact on A/B/C, 69.6% on D). Patches are minimal (mean 2.1 lines, max 6).

RQ3: Does the method survive a criterion-reversal audit? (Methodology lesson)

The original repair criterion (prove/k-induction) misjudged 78 correct repairs on axi/uart (golden itself UNKNOWN under prove). BMC reverification: 303/303 pass — repair rate 74.3% → 100%.

Lesson: repair criterion must be validated against golden first (criterion-vs-golden consistency check).

RQ4: Cost & performance

- Total LLM cost: \$11.22 for the four-setting main experiment (408 runs, \$0.0275 avg); full ledger \$14.36 (839 calls).
- Per-setting (all n=102): A \$2.78 / B \$2.72 / C \$4.17 / D \$1.56; tokens A 1.99M / B 1.76M / C 3.50M / D 1.07M.
- C-window compression: main dataset traces are short (median 6 cycles); window=8 yields only 4.1% aggregate reduction — semanticization cost is dominated by summary text, not trace length.
- Verification: BMC 8-concurrency Gate-2 reverify = 68 jobs / 3.4 min; axi golden ~87s, uart_rx ~151s.

Open items before paper submission

- ☒ Fill D setting numbers (102 runs, \$1.56, loc 49.0%, repair 100%)
- ☒ FVDebug quantitative comparison via D (D cheap but lower precision than B; B remains main setting)
- ☐ Human interpretability rating (small study) for C/D
- ☒ Paper sections: Intro / Method (2026-08-04); Related Work + Threats to Validity / Conclusion (2026-08-04)
- ☐ Full-draft integration (paper/main.md) + cross-section consistency pass

4. Related Work

Scope

PreCex sits at the intersection of four lines of work: (i) formal-verification failure analysis and counterexample understanding; (ii) LLM-based RTL debugging and repair; (iii) assertion genera-

tion and verification sufficiency; and (iv) RTL debugging benchmarks. This section positions our pipeline against each, with emphasis on the evidence-representation axis that PreCex ablates.

Counterexample-driven debugging and root-cause analysis

FVDebug [arXiv:2510.15906] is the closest competitor. It builds a causal DAG from a failing formal trace, scans suspicious nodes with batched LLM calls, produces a root-cause narrative through an agent, and generates fixes, demonstrated on two industrial counterexamples. PreCex shares the counterexample-understanding core: our Setting D is a deterministic, fully open-source realization of the FVDebug-style causal graph (failed assertion + fault cone + full-cycle state trace + trigger condition). FVDebug leaves three axes unaddressed that PreCex controls for: it depends on a commercial toolchain, it reports no controlled ablation of the evidence representation, and it does not quantify verifier sufficiency. PreCex is fully open-source (iverilog/yosys/SymbiYosys/Z3), compares four evidence representations on the same 34-sample dataset, and audits the verifier with mutation analysis, vacuity control, and an independent T2 review agent.

GROVE [arXiv:2511.17833] organizes assertion-failure debugging knowledge as a layered knowledge tree, validates knowledge items with model checking at training time, and performs budget-aware retrieval at test time, consistently improving pass@1/pass@5. GROVE and PreCex are complementary: GROVE accumulates reusable knowledge across samples, whereas PreCex consumes the single failing counterexample in a closed localize-repair-reverify loop without a history store. A GROVE-style knowledge layer is a natural extension of our pipeline.

Open-source LLM-driven formal verification [arXiv:2607.28877] uses the same open toolchain (Yosys/SymbiYosys/Z3) with LLM-guided counterexample-driven iterative repair, stopping on k-induction proof or budget exhaustion, and reliably fixed only 1 of 6 benchmarks. This work validates the open-source pipeline as feasible and catalogs its failure modes. PreCex differs in two ways: it adds an evidence-semanticization layer (that work feeds raw counterexamples directly to the LLM) and it evaluates on a larger, controlled L3 dataset with a golden-first verification criterion.

FormalRTL uses a software reference model as a formal specification and closes a plan/synthesis/equivalence-checking loop, including a counterexample simplifier that highlights only the failing function and mismatched signals. PreCex makes no reference-model assumption: it works from assertions plus a counterexample alone, which is the common pre-synthesis setting where no golden model exists.

ChipAgents RCA (commercial, 2026) combines a waveform-understanding engine with a prover-verifier agent pair and self-consistency ranking, and reports localization of a three-cycle race condition in industrial deployment. It corroborates that cross-cycle root-cause analysis is a real industrial pain point, but is a closed black box with no ablation detail.

LLM-based RTL repair

VeriPilot [arXiv:2606.23759] uses a golden reference model with CDFG signal tracking and raises GPT-4o repair rates on CVDP from 54.3% to 85.71%. **VeriDebug** [arXiv:2504.19099] learns a contrastive embedding and guides correction, reaching 64.7% Acc@1 on a combined localization-plus-type task. **RTLFixer** applies retrieval-augmented repair to syntax errors. These works are not counterexample-driven: they rely on reference models, fine-tuned embeddings, or syntax-only fixes. PreCex deliberately targets the pre-synthesis setting with no reference model and treats the evidence representation rather than model capacity as the studied variable.

Fixbench-RTL and **HDL-FixBench** are repair benchmarks spanning syntax/functional/security

domains (Fixbench) and repository-scale instances (57 instances from OpenTitan/CVA6/Ibex, best model 40.3%; rejected by ICLR 2026). HDL-FixBench’s rejection highlights pitfalls we engineered against: small scale, insufficient diversity, disputed patch thresholds, missing non-LLM baselines, and contamination risk. BugBench-PS addresses these with 34 L3 samples across 6 modules and 7 error classes, per-sample golden double-checking, tool-execution adjudication, and a documented contamination statement.

Assertion generation and verification sufficiency

AssertLLM [ASPDAC’25] and **AssertLLM2** generate SVA from natural-language specifications plus waveforms; AssertLLM2 introduces an 83-design assertion benchmark and evaluates assertions through mutation-based bug detection in a bug-hunting setting. **AssertGen** [ATS’25] extracts verification goals via chain-of-thought and bridges cross-layer signals, with mutation-testing coverage evaluation. **LASSO** [MLCAD’25] generates safety properties with explicit vacuity checking and coverage feedback, detecting 5 real bugs in OpenTitan. **LintLLM** [GLSVLSI’25] performs LLM linting with mutation-based defect injection.

These works generate or audit assertions; PreCex consumes assertions plus counterexamples to localize and repair. Their evaluation practices (mutation coverage, vacuity, bug-hunting) provide the methodological precedent for our sufficiency audits: 88.5% strong mutants and 81.8% constant mutants killed, with delete-mutation vacuity control.

Benchmarks and evaluation methodology

Tool-execution adjudication is an accepted evaluation practice in the EDA-AI literature: Rule2DRC (ICML 2026) compares 13,921 layout-rule results and PostEDA-Bench evaluates 145 DRC/PPA tasks by executing design tools. PreCex follows the same principle: repair success is adjudicated by compilation, weak-testbench regression, and formal BMC, never by the LLM’s self-report. Engineering-semantics datasets (OpenRTLSet, EDA-Schema-V2, ChipLingo) informed our evidence schema design.

Positioning

Work	Evidence representation	Open-source	Cross-cycle (L3) focus	Controlled evidence ablation	Sufficiency quantification
FVDebug	causal graph	no	partial	no	no
GROVE	knowledge tree	no	no	no	no
2607.28877	raw cex	yes	partial	no	no
FormalRTL	cex simplification + reference model	partial	no	no	no
VeriPilot	CDFG + reference model	no	no	no	no

Work	Evidence representation	Open-source	Cross-cycle (L3) focus	Controlled evidence ablation	Sufficiency quantification
VeriDebug	learned embedding	no	no	no	no
AssertLLM2	— (assertion gen.)	—	—	—	mutation-based
PreCex	raw / structured / semantitized / causal graph	yes	yes (34 samples)	yes (A/B/C/D)	yes (mutation/vacuity/T2)

Our three differentiators are each evidenced by controlled experiments in this paper: (1) a fully open-source, reproducible counterexample-driven pipeline; (2) a four-setting evidence-representation ablation with paired significance testing; (3) a cross-cycle (L3) benchmark with a verifier-sufficiency loop (mutation, vacuity, and an independent T2 audit). In addition, the criterion-reversal audit (Results, RQ3) contributes a methodological lesson: a repair criterion must be validated against the golden design before it is used to judge repairs.

5. Threats to Validity and Conclusion

Threats to validity

Internal validity

Repair criterion. All reported results use BMC (verify.sby) as the primary criterion, consistent with golden verification. We found that prove/k-induction does not converge on gated sequential assertions in the axi/uart modules: the golden design itself returns UNKNOWN, and the original protocol misjudged 78 correct repairs as failures. A golden-first rule now guards the criterion — any candidate criterion must first pass on the golden design. Residual risk remains that a defect manifests only beyond the probed BMC depth; we probe to bug depth + 2 and cross-check against golden verification, but this is not a completeness guarantee.

LLM nondeterminism. We run 3 seeds per sample per setting and report paired statistics. The pipeline itself is deterministic (parsing, extraction, verification); the only stochastic component is the LLM inference.

Repair-success definition. Success requires compilation, weak-testbench regression, BMC with no counterexample, and golden double-check. A patch that satisfies the probed assertions while changing behavior beyond the property set would pass; the T2 audit bounds this risk by checking for interface changes and assertion tampering, but it is not a full equivalence check.

External validity

Scale and scope. The benchmark contains 34 L3 samples across 6 modules (fifo_sync, uart_tx, uart_rx, axi_lite_slave, fsm_ctrl, counter_alu) and 7 error classes. The modules are open-source

teaching-grade RTL; we do not claim industrial protocol coverage. Handshake and state-transition defects remain the hardest classes to locate (loc_top1 27.8% and 36.1% respectively), and uart_rx’s baud-timing constant mutations are weakly killed (const 40.9%) — a documented limitation of structural assertions for timing-constant shifts.

Contamination. Common modules may appear in LLM training corpora. We mitigate with construction-date records, per-sample provenance, and error-class diversity, but residual memorization cannot be fully excluded.

Single LLM provider. All 408 runs use MiniMax M3. Absolute rates may shift with model versions or providers; the relative evidence-representation effect is the claim we report.

Construct validity

loc_top1 records whether the model’s top-ranked candidate line is inside the injected defect’s reference diff; repair rate is measured under the BMC criterion; cost is metered by API tokens. These constructs match the repair-benchmark literature. Statistical claims use paired McNemar tests on per-sample localization outcomes (102 pairs per comparison): B vs A $p=0.0035$ and B vs D $p=0.0164$ are significant; B vs C $p=0.404$ and C vs D $p=0.186$ are not. All comparisons are reported, including non-significant ones.

Conclusion

We presented PreCex, an open-source, counterexample-driven pipeline for pre-synthesis RTL defect localization and repair, and studied the evidence-representation axis that prior counterexample-understanding systems leave uncontrolled. On a 34-sample cross-cycle (L3) benchmark with 408 LLM repair runs:

- All four evidence settings (raw logs, structured evidence, semanticized counterexamples, FVDebug-style causal graphs (D)) reach 100% repair under a consistent BMC criterion; the representation determines how precisely and how cheaply, not whether.
- Structured evidence (B) achieves the highest localization precision (61.8%), statistically significant over raw logs (47.1%, $p=0.0035$) and causal graphs (49.0%, $p=0.0164$). Causal graphs (D) are the cheapest (\$1.56 vs \$2.72, 43% lower); semanticization (C) adds cost without significant precision gain.
- Repairs are trustworthy by construction: mutation analysis kills 88.5% of strong and 81.8% of constant mutants, vacuity control passes, and the independent T2 audit passes 408/408 runs with zero interface or assertion-tampering changes and zero median localization deviation.
- The criterion-reversal audit (BMC over prove/k-induction, with a golden-first rule) is a methodological lesson for automated repair pipelines: the evaluation criterion itself must be validated before repairs are judged.

Future work includes scaling to industrial protocols and larger assertion sets, stronger mutation families, multi-provider and multi-model evaluation, a GROVE-style knowledge layer for cross-sample reuse, and reinforcement learning with formal proofs as reward.